

A Lean Audit of the IUT III Corollary 3.12 Corridor

IUT Lean formalization project

June 1, 2026

Abstract

We report on a first Lean formalization of the logical corridor from IUT III, Theorem 3.11 to Corollary 3.12. The formalization is deliberately neutral with respect to the correctness of IUT as a proof of the abc conjecture. It isolates the ordered-real comparison required at the end of Corollary 3.12, tracks the finite q^{j^2} and label-quotient phenomena that enter the public dispute, and records exactly which source-facing construction hypotheses are still external.

The current code proves conditional theorems. If the hull/log-volume upper-ray comparison supplies $-|\log(q)| \leq -|\log(\Theta)|$, and if the displayed hypothesis $-|\log(\Theta)| \leq C_\Theta |\log(q)|$ is supplied with $|\log(q)| > 0$, then Lean proves $C_\Theta \geq -1$. If the upper-ray membership is strict, Lean proves $C_\Theta > -1$. In the boundary case $C_\Theta = -1$, Lean proves equality of the q -pilot reading and the Θ -hull boundary, and then classifies the local Frobenioid p_v^N -shift behavior under the actual upper-ray quotient. This does not prove Corollary 3.12 from Mochizuki’s published hypotheses; it identifies the remaining mathematical construction problem.

After a source audit against IUT I–IV, Mochizuki’s 2026 LeanForm report, and Scholze–Stix, the status is unchanged in kind but sharper in formulation: the final ordered-real implication is formalized, and several collapse-prone comparison levels are separated, but the code still assumes the source construction of the multiradial/HDD/SHE/IPL data that IUT III assigns to Theorem 3.11 and its reorganization.

1 Aim and Sources

The project formalizes a narrow and contested part of Mochizuki’s inter-universal Teichmüller theory. The source base is the local copy of IUT I–IV, Mochizuki’s 2026 report on formalization, and the Scholze–Stix note. The current target is the portion that Mochizuki’s report calls Stage 1:

$$\text{IUT III, Theorem 3.11} \xrightarrow{\text{APT+IPL}} \text{“3.11.5”} \xrightarrow{\text{SHE+IPL}} \text{IUT III, Corollary 3.12.}$$

The notation “3.11.5” is not a statement in IUT III. It is Mochizuki’s 2026 name for the reorganized boundary where the hull-plus-determinant operation is moved before the final simultaneous comparison of the Θ -pilot and the q -pilot.

The formalization is source-facing but not source-complete. It does not construct Hodge theaters, Frobenioids, log-shells, prime strips, or theta links from first principles. Instead, these are represented by typed records with named hypotheses. The point of the current work is to test whether the claimed formal corridor is coherent once every comparison level, label quotient, ordered real line, hull boundary, and quotient collapse is made explicit.

2 Framework and Tooling

The repository is organized as a Lean verification project, not as a prose annotation project. The source papers in `docs/` are treated as the mathematical source of truth. The formalization then

moves through three layers:



The paper draft records only the current mathematical state of the formalization. It is not a research log.

The Lean code uses typed interfaces to prevent comparison-level drift. A datum that is merely a representative j^2 -coordinate computation has a different type from a balanced full-label datum, and both differ from the final hull/log-volume comparison datum. Likewise, ordered real-line transports, Hodge/SHE preservation records, upper-ray membership, and quotient-collapse audits are separate structures. This is the main anti-collapse mechanism in the framework: a proof cannot silently use a raw coordinate equality where the Corollary 3.12 corridor asks for a hull/log-volume comparison.

The module roles are:

- **IUTStage1SourceCore**: core real-line, pilot, hull, quotient, and C_Θ algebra.
- **IUTStage1Remark312Absorption**: Remark 3.12.2(v) absorption interfaces.
- **IUTStage1Source**: finite labels, Gaussian data, SHE/Hodge interfaces, and route theorems.
- **IUTStage1Experiments**: compiled endpoint checks and diagnostic theorems.

The experiment file is not a substitute for proof. It is a curated surface that forces the important route theorems to elaborate together and exposes Boolean/report endpoints such as whether the present stage settles the dispute. At present that endpoint is deliberately false.

The build discipline is simple. Each mathematical increment is checked by Lean via `lake build Iut.Stage1.IUTStage1Experiments`; the paper is rebuilt from this LaTeX source; and the code is scanned for explicit proof placeholders such as `sorry`, `admit`, and `axiom`. The long-term collaboration plan is compatible with Apodixis: the current typed route endpoints and missing-obligation records give a natural surface for a collaboration tool to show which mathematical interfaces are proved and which remain open.

3 Dispute Interface

In IUT III, Step (x) treats the procession-normalized log-volume as unchanged under (Ind1) and (Ind2), while (Ind3) supplies a one-sided upper-semi comparison. Step (xi) then passes through multiradiality, IPL, SHE, holomorphic hulls, determinants, and log-volume to place the q -pilot reading in the signed-log upper ray determined by the Θ -pilot:

$$-|\log(q)| \in \mathbb{R}_{\leq -|\log(\Theta)|}.$$

The active source-facing Corollary 3.12 route does not use equality of possible-image regions across an (Ind3) generator. Such an equality exists only in separate quotient bookkeeping interfaces as an explicit stronger hypothesis; the route uses the one-sided upper-semi corridor as its Step (x) input. The current Theorem 3.11 source record now bundles structured IPL/SHE/APT input, coric multiradial output, possible theta-pilot images, the 0-column/1-column pilot distinction, distinct

log-shell treatments, and the Hodge-theater history guard. This is still a source interface, not a construction of Theorem 3.11. The IPL layer has been separated into a typed log-volume transport datum: it carries the source and target Hodge-theater sides, identifies the prime-strip link endpoints with the input and output prime strips, records equality of the transported real log-volume, and retains the no-history-identification guard. It is not yet a construction of the full log-link. The current SHE layer now has a synchronization constructor: equality of the canonical one-label Hodge–Arakelov theta-value degree constructs the factored square/full-label preservation package used by the route. This is a Stage 1 source constructor for the available finite-label shadow of SHE, not yet the full simultaneous-holomorphic-expressibility theorem. The hull/determinant layer now has a source constructor attached to the Theorem 3.11 record. It reads the algorithmic certificate and SHE alignment from that record, keeps the record’s theta possible-image union tied to the hull endpoint, and derives the determinant-bound comparison $q_{\text{signed}} \leq \theta_{\text{signed}}$. The bridge used by this source data can now be constructed from a displayed common hull containing the target union and a determinant/log-volume bound for that hull; the analytic construction of the hull and determinant remains external source data. The strongest Hodge–Arakelov packet-target route no longer takes the target-side log-volume equality or q -pilot positivity as independent inputs: target equality is derived from the SHE synchronization constructor, and positivity is read from the hull/determinant constructor. The source-side chart calibration is no longer a raw equality input: it is a full-label theta/Hodge calibration certificate, from which equality of the cusp log-volume compatibility records is derived extensionally. The final numerical step is elementary:

$$-|\log(q)| \leq -|\log(\Theta)| \leq C_{\Theta} |\log(q)|, \quad |\log(q)| > 0 \implies C_{\Theta} \geq -1.$$

Scholze–Stix object to the comparison because the Θ -side data carries j^2 -scaled degrees, while a naive identification of ordered real lines appears to remove the room needed for these scalars. The formalization therefore separates four comparison levels:

representative j^2 data
balanced sign-compatible data
averaged cusp-label data
hull/log-volume data.

The final route consumes only the hull/log-volume comparison level. Lean rejects the balanced sign-compatible level as sufficient for the final comparison.

4 Source Audit

The local sources impose the following constraints on the formalization. IUT I constructs the initial Θ -data and the $\Theta^{\pm\text{ell}}$ NF-Hodge theaters. It is not merely a notation layer: the paper’s central point is that the Hodge theaters separate the additive and multiplicative combinatorial dimensions of a number field. IUT II supplies the Hodge–Arakelov theta evaluation, theta values at ℓ -torsion points, Gaussian Frobenioids, conjugate synchronization, and the multiradial/coric phenomena needed later. The current code models only finite label and Gaussian degree shadows of this material.

IUT III is the direct source for the formal corridor. Step (x) says that procession-normalized log-volumes are invariant under (Ind1) and (Ind2), while (Ind3) becomes an upper-semi inequality after applying log-volume. Step (xi) uses the multiradial construction, IPL, SHE, holomorphic hulls, determinants, and the 1-column log-Kummer correspondence to put the q -pilot log-volume in the signed-log upper ray $\mathbb{R}_{\leq -|\log(\Theta)|}$. Step (xii) warns that replacing global realified Frobenioids by

local Frobenioids loses calibration because local objects admit p_v^N -type endomorphisms. Remark 3.12.2 further stresses that the argument depends on weakened data, distinct q/Θ labels, upper-semi indeterminacy, and the asymmetry between the 0-column Θ -pilot and the 1-column q -pilot. The current Step (x) packet source reflects this asymmetry by packaging the 1-column q -pilot non-interference, Kummer/forgetting product preservation, source packet calibration, and packet-normalized target calibration in a single realified-Frobenioid log-Kummer source object.

IUT IV uses the resulting log-volume estimate as input to elementary log-volume and Diophantine inequalities. Our IUT IV code is therefore only conditional algebra at present: it records how later estimates consume a Corollary 3.12-shaped bound, but it does not prove that bound from the full arithmetic constructions.

Mochizuki’s 2026 report explicitly recommends Stage 1 as the implication *IUT* III, Theorem 3.11 $\Rightarrow$ Corollary 3.12, with a reorganized intermediate “3.11.5” obtained by moving the hull-plus-determinant operation before the final simultaneous comparison. The same report states that the current official code is skeletal. Our formalization is best read as a more diagnostic Stage 1 interface: it does not replace Stage 2, because it has not constructed Theorem 3.11.

Scholze–Stix argue that, after simplification, the comparison loses the j^2 -scaled theta contribution and becomes essentially contentless. The Lean code therefore treats the j^2 -representative level, sign quotient, averages, square/full-label preservation, ordered real-line transports, and hull/log-volume endpoint as different types. This confirms that the final Corollary 3.12 endpoint cannot be obtained from the balanced sign-compatible level alone. It does not prove that Mochizuki’s source constructions supply the missing higher-level data.

5 Formalized Corridor

The relevant implementation lives in four Lean modules:

- `Iut/Stage1/IUTStage1SourceCore.lean`
- `Iut/Stage1/IUTStage1Remark312Absorption.lean`
- `Iut/Stage1/IUTStage1Source.lean`
- `Iut/Stage1/IUTStage1Experiments.lean`

5.1 Finite label model

The code defines a finite $\mathbb{F}_\ell$ -label model, the quotient $|F_\ell|$ consisting of the zero class together with $\mathbb{F}_\ell^\times/\{\pm 1\}$, and the canonical half-range representatives

$$j = 0, \dots, (\ell - 1)/2.$$

The Stage 1 label model is now connected to the bad-local theta-root data of IUT I, Definition 3.1(e),(f). A new source package consumes the quotient- $\mathbb{Z}$ data of a bad local theta-root model and derives the Stage 1 full label from its canonical coordinate. It also carries the canonical-generator-up-to-sign witness. In the concrete $\mathbb{Z}/\ell\mathbb{Z}$ quotient case this derived full label is exactly the nonzero canonical full label corresponding to coordinate 1. It proves, for the descended full-label degree model:

$$\deg(\Theta_j) = j^2 \deg(q)$$

for the degree-level Gaussian endpoint, with sign invariance on the full-label quotient. It also proves that $j = 1$ and $j = 2$ give distinct degrees exactly when $\deg(q) \neq 0$. Thus the q^{j^2} phenomenon is present in the model; it is not hidden inside a scalar identification.

The raw coordinate square-weight profile on $\mathbb{Z}/\ell\mathbb{Z}$ is kept as a diagnostic object and is not identified with the descended $|F_\ell|$ theta exponent. The coordinate average, the nonzero-coordinate average, and the absolute-label average are all formalized separately. Lean proves the exact denominator relations and shows that these averages differ strictly for nonzero environment degree. This prevents silently replacing the paper's ℓ -denominator convention by an auxiliary nonzero-only average.

5.2 Indeterminacy and upper-semi route

The Step (x) corridor is represented as:

$$A_0 = A_1 = A_2 \leq U_{\text{Ind3}},$$

where A_0, A_1, A_2 are the before-, after-(Ind1)-, and after-(Ind2)-averaged log-volumes, and U_{Ind3} is the (Ind3) upper-semi target. From this data and the hull/determinant identifications, the code constructs the Step (xi) upper-ray datum:

$$q_{\text{avg}} \leq \Theta_{\text{hull}}.$$

The same route constructs the two q -pilot readings in Step (xi-g) and proves that they are equal to the same fixed real value.

5.3 SHE and Hodge-descent boundary

The current SHE layer is conditional. It has typed records for:

structured SHE square-weight transport
factored square/full-label preservation
Hodge-descent packet transport
nonarchimedean (Ind3) source-target alignment.

Given these inputs, Lean constructs the weighted theta comparison consumed by the final route:

$$q_{\text{signed}} \leq \Theta_{\text{signed}}.$$

The nonarchimedean (Ind3) entry, factored square/full-label preservation, and Hodge-descent packet transport now compose directly to this comparison. With the q -pilot positivity condition and the displayed C_Θ -bound, the same composed route instantiates the signed C_Θ endpoint and proves $C_\Theta \geq -1$. The instantiated endpoint also exposes the strict and boundary alternatives: strict source comparison gives $C_\Theta > -1$, while $C_\Theta = -1$ forces equality of the signed q - and Θ -readings and makes the Θ -reading negative. These facts are now also packaged as a single route-level alternative: boundary equality with negative Θ -reading, or strict $C_\Theta > -1$. The code also proves a Gaussian-to-factored-SHE bridge at the present finite degree level. In particular, equality of the source and target environment degrees gives full-label value preservation; nonzero-label target bounds are enough because the zero label has square weight zero. The newest route uses this bridge before applying the nonarchimedean endpoint: Gaussian degree evaluations with preserved environment degree now construct the factored square/full-label preservation package consumed by the

nonarchimedean C_Θ dichotomy. This removes one primitive preservation input from that route, but it still assumes the nonarchimedean (Ind3) alignment and the Hodge-descent/SHE source audit. There is also an identity-coordinate specialization. In that case equality of Gaussian degree at the canonical full label 1 implies preservation of the environment degree, so the nonarchimedean C_Θ dichotomy no longer takes environment-degree preservation as an independent hypothesis. The Hodge–Arakelov/Gaussian bridge is now source-facing at the current degree level. A theta-value evaluation source carries the bad-local theta-root label data and a theta-monoid degree d , and constructs the Gaussian degree evaluation

$$\deg_{\text{gau}}(\lambda) = e(\lambda) d,$$

where $e(\lambda)$ is the absolute-label theta exponent. The zero label has degree 0, and the canonical label 1 has degree d . Thus equality at the canonical label derives equality of the environment degree and feeds the factored SHE square/full-label obligations. The raw nonarchimedean entry alignment has also been refactored through a source-facing nonarchimedean upper-semi object equipped with 1-column q-pilot log-Kummer data. The upper-semi inequality remains the Step (x) input; the 1-column datum records the q-pilot column and log-Kummer non-interference before converting to the raw real equalities consumed by the current route. For the source side of this compatibility, the code can now use the existing finite tensor-packet chain: Kummer transfer holomorphicF $\rightarrow$ holomorphicD and mono-analytic forgetting holomorphicD $\rightarrow D^+$ preserve product log-volume. Thus the source-side upper-semi equality may be derived from this chain rather than postulated as a bare equality. The newest source package attaches this finite Kummer/forgetting chain to the bad-local theta-root cusp-label source. It carries the canonical nonzero $|F_\ell|$ -label and derives the nonarchimedean (Ind3) source equality from product-log-volume preservation along

$$\text{holomorphicF} \rightarrow \text{holomorphicD} \rightarrow D^+.$$

The packet local object, local entry source, mono-analytic product, and holomorphic source product are now packaged as one source-side log-Kummer packet calibration object. The object itself is still supplied source data, but the equality from entry source to (Ind3) source is derived from it together with the Kummer and forgetting maps. The source calibration can now also be constructed from the explicit packet-local/source alignment, the packet-local (Ind3)-source alignment, and the two product-preservation maps. Thus the route no longer has to accept the holomorphic source-product equality as a primitive field when these Step (x) alignments are available. The source and target packet calibrations are now joined by a finite log-Kummer packet-correspondence source carrying the q-pilot column and non-interference guard. Its target face carries the common packet-normalized capsule value, and Lean derives the direct theta-to-entry, theta-to-packet, and entry-to-(Ind3)-target equalities consumed by the upper-semi route. The target calibration no longer contains an independent real representative for this common value: all target equalities are stated directly against the packet-normalized capsule log-volume. This removes one artificial identification layer from the Step (x) interface. The local entry target can now also be constructed from the packet-normalized capsule value itself. A packet-normalized entry-target source builds a nonarchimedean upper-semi entry whose target finite log-volume is definitionally the packet-normalized capsule value. The corresponding target calibration therefore no longer needs an independent proof for the entry-target equality; only the theta-average and (Ind3)-target identifications with the same packet-normalized value remain as inputs at this level. The latter can now be reduced further: if the Step (x) target-side (Ind3) alignment gives equality of the theta average with the (Ind3) target, then Lean derives the (Ind3)-target-to-packet-normalized equality from this alignment and the theta-average-to-packet-normalized equality. The packet-correspondence source itself can now be

built from this reduced target data together with the source packet/product calibration. Thus the correspondence constructor no longer needs a completed target-calibration record when the target entry and theta-target alignment are available. It can also be built directly from the packet-local source alignments and the same target theta alignment. In this form both the source packet calibration and the target packet calibration are generated before the q-pilot non-interference endpoint is projected. The same reduction has been lifted through the realified-Frobenioid packet source: realified Kummer compatibility still supplies the $\text{holomorphicF} \rightarrow \text{holomorphicD} \rightarrow D^+$ transfer, while the target face is now generated from the packet-normalized entry target and theta-target alignment. The realified source can now also generate the source packet calibration from packet-local/source alignments before forming the correspondence; consequently the strongest realified Step (x) constructor no longer needs a completed source calibration as input. The first-pass Gaussian route also has a one-correspondence form: it consumes a finite log-Kummer packet-correspondence source and projects the source packet/product calibration and target packet-normalized calibration from that single object before entering the upper-semi corridor. The current strongest Step (x) entry source places this correspondence below the realified Frobenioid layer: the selected nonarchimedean upper-semi entry and the realified-Frobenioid packet source are bundled together, and Lean projects the Kummer transfer, mono-analytic forgetting transfer, source calibration, target packet calibration, and q-pilot non-interference guard from that single source object. There is now a reduced constructor for this entry source: entry membership, packet-local/source alignment, packet-local (Ind3)-source alignment, packet-normalized target data, and theta-target alignment construct the realified packet source internally and then derive the upper-semi inequality. The bad-local theta-root label source is now attached to this realified entry source. The corresponding Hodge–Arakelov route consumes one object carrying the canonical nonzero theta-root label, the selected upper-semi entry, and the realified-Frobenioid log-Kummer packet source before entering the C_Θ -dichotomy. This theta-root attached source now has the same reduced constructor: the theta-root label package is combined with entry membership and the packet-local/target alignments, and Lean projects the canonical-label facts together with the Step (x) upper-semi endpoint. The Hodge/SHE/IPL/hull route has the same realified endpoint: its alignment record supplies the Hodge–Arakelov Gaussian comparison, canonical-one preservation, q-pilot positivity, and theta/Hodge log-volume calibration, while the Step (x) data enters through the single theta-root attached realified source. The route now also has a source-reduced endpoint: the theta-root label package, entry membership, packet-local source alignments, packet-normalized target alignment, and realified Frobenioid compatibility construct the Step (x) source inside the theorem before proving the same $C_\Theta = -1$ boundary alternative or strict $C_\Theta > -1$ conclusion. The Milestone 2 hull layer now contains a direct formalization of Remark 3.9.5(iii)’s approximant family. For a region P , an approximant record carries a hull H , proves $H \subseteq \phi(P)$, and proves

$$\mu_{\log}(P) \leq \mu_{\log}(H) \leq \mu_{\log}(\phi(P)).$$

The canonical hull $\phi(P)$ is itself an approximant, and every approximant is fixed by hull formation. This is the compactness guard emphasized in the paper: the indeterminacy in choosing a hull approximant remains inside $\phi(P)$. The bounded-family version of Remark 3.9.5(v)–(vi) is now represented by a family-hull quotient source. For a family P_β , the code forms the hull of the union, uses it as the collapsed subset in the upper-semi quotient, and proves that any two nonempty possible regions have the same quotient image. The Corollary 3.12 possible-image approximant source projects to this family-hull quotient source, so the possible-image family used in the upper-ray route also carries the corresponding choice-independence theorem. Since the q-pilot region is contained in the chosen approximant and the approximant lies in the same family hull, the code also proves that a nonempty q-pilot region has the same family-hull quotient image as any nonempty

possible theta-image choice. The same approximant now feeds the Step (xi) upper-ray constructor. If the q -pilot region lies in an approximant H , and if $\mu_{\log}(H)$ is the determinant-normalized log-volume, then Lean derives

$$\mu_{\log}(q) \leq \mu_{\log}(H) \leq \mu_{\log}(\phi(P)), \quad \mu_{\log}(q) \leq \det_{\text{norm}}.$$

Thus this part of the corridor no longer has to use only the canonical hull itself; it can use any certified hull-approximant inside $\phi(P)$. This has now been specialized to the possible-image union in Corollary 3.12: one source object carries the family of possible theta images, the union P , a certified approximant $H \in \Phi(P)$, q -pilot containment $q \subseteq H$, and determinant normalization of $\mu_{\log}(H)$. Lean then derives the whole chain

$$\mu_{\log}(P) \leq \mu_{\log}(H) \leq \mu_{\log}(\phi(P)), \quad \mu_{\log}(q) \leq \mu_{\log}(H) \leq \det_{\text{norm}}.$$

The determinant side has also been sharpened at finite log-volume level. Following Remark 3.9.5(vii), (Ob3-1)–(Ob3-3), a localization summand now carries a rank > 1 , a bundle log-volume, a structure-sheaf adjustment, and a positive weight. A determinant source forms the finite weighted sum of the adjusted summands, then derives the positive tensor-power log-volume and the normalized determinant log-volume by definition. This replaces supplied tensor-power and normalized values for this new route, but it remains a log-volume skeleton rather than a construction of arithmetic vector bundles. The localization input itself has been refined one step further: a finite direct-summand source computes the bundle log-volume as a sum of direct summand log-volumes before projecting to the localization summand interface. Thus the weighted determinant can now be built from finite direct-summand localization sources. The Ob3-1-2 structure-sheaf adjustment is also named explicitly. A structure-sheaf adjusted source carries

$$L_{\text{adj}} = \sum_{\gamma} L_{\gamma} - L_{\mathcal{O}(-)}$$

and proves that its weighted value is the same adjusted summand consumed by the determinant source. Thus the current determinant route separates direct summands, subtraction of the $\mathcal{O}(-)$ contribution, positive weighting, tensor-power formation, and normalization as distinct Lean interfaces. The possible-image approximant route can now consume this weighted determinant source directly. The equality between the approximant log-volume $\mu_{\log}(H)$ and the weighted normalized determinant value is now carried by a named Ob3-3 compatibility certificate. After this certificate is supplied, the old determinant record is only a projection used by the common upper-ray interface. At the package-ledger boundary, the hull/determinant source data can now be constructed from a concrete common hull K , a proof $P_{\text{out}} \subseteq K$, and a determinant/log-volume estimate $\mu_{\log}(K) \leq \theta_{\text{signed}}$. The resulting Lean endpoint returns target-union containment, the determinant bound, all-target bounds, and equality between the package's named hull/determinant bridge and this constructed bridge. Thus the remaining source obligation is the construction of K and its bound, not an opaque common-bound bridge. Remark 3.9.5(vii), (Ob4), is represented at the same finite level: a naive Frobenius/tensor-power log-volume object records the passage $L \mapsto ML$ for $M > 0$, and Lean proves that normalized log-volume recovers L . The weighted determinant source projects to this object, so the M -th tensor-power copy does not add a further log-volume ambiguity in the current corridor. A comparison object also proves that two positive tensor-power presentations with the same base log-volume have the same normalized log-volume, even when the tensor degrees differ. Below the packet layer, the code now has a first realified Frobenioid degree source. A finite Frobenioid object carries a divisor degree, a unit log-volume, and their realified log-volume; a tensor-packet product

may be declared to have product log-volume equal to this realified value. A finite tensor-product operation on these degree objects adds divisor degrees and unit log-volumes, and Lean proves that the realified log-volume of the tensor product is the sum of the two realified log-volumes. This operation is also lifted to the tensor-packet source layer: if an existing tensor-packet product source is certified to carry the tensor-product Frobenioid degree, Lean derives additivity of its packet product log-volume. A Frobenioid Kummer-compatibility certificate then projects to the older coric tensor-packet transfer by deriving equality of product log-volumes from equality of realified Frobenioid log-volumes, while retaining the Hodge-theater history guard. This is still a finite degree-level shadow, not a construction of Mochizuki's Frobenioids, but it moves one equality source below the packet-correspondence interface. The global-to-local realified restriction from IUT I, Example 3.5 is also encoded at degree level: the restricted global prime log-volume is $[K_v : (F_{\text{mod}})_v]^{-1}$ times the local prime log-volume, and Lean proves that multiplying by the extension degree recovers the local prime log-volume. The same restricted value now normalizes the local p_v^N -shift source: the local shift is expressed with this restricted global prime step, and the automorphism-or-zero-step criterion is restated in terms of the restricted realified prime log-volume. IUT I, Example 5.1 and Remark 5.1.1 distinguish the diagonal Frobenioid $F_{\langle J \rangle}$ from the product of underlying categories F_J : the former carries constant distributions in j , while the latter allows arbitrary distributions. The code now models this finite realified log-volume distinction. A diagonal distribution has total log-volume $|J|$ times its common value, with the realified restriction factor propagated through this sum; a distribution with two unequal coordinates is proved not to come from a diagonal constant distribution. IUT I, Example 5.4(vii) and Remark 5.4.2 are represented by a finite pivotal distribution object: a chosen pivotal index supplies the common diagonal log-volume, the summed interpretation gives total $|J|$ times this pivotal value, and the averaged interpretation recovers the pivotal value itself. IUT I, Definition 5.5 is now represented at finite log-volume level by an NF-bridge mode object: the bridge target is either the summed capsule log-volume or the normalized averaged capsule log-volume. For the same capsule distribution, Lean proves that the summed bridge target is $|J|$ times the averaged bridge target; for a pivotal distribution, these specialize to $|J|$ times the pivotal value and to the pivotal value itself. The adjacent Θ -bridge and Θ NF-gluing condition of Definition 5.5(iii) are formalized at the same finite level: the Θ side reads the common capsule total, and a compatibility record requires the NF- and Θ -bridge morphisms to use the same index bijection and the same capsule distribution. Lean then proves that a summed NF target equals the Θ -target, while an averaged NF target becomes the Θ -target after multiplication by $|J|$. Corollary 5.6(iii) is represented by a finite gluing torsor: once compatible NF and Θ bridge halves have been fixed, the gluing parameter is modeled by the additive F_ℓ -torsor on $\mathbb{Z}/\ell\mathbb{Z}$. Lean proves zero action, additivity of the action, existence and uniqueness of the translation between two gluing parameters, and preservation of the shared index-bijection and capsule data. This layer is now connected to the nonarchimedean packet correspondence: a realified-Frobenioid log-Kummer packet source projects the Frobenioid compatibility data to the Kummer $F \rightarrow D$ transfer and to the mono-analytic forgetting transfer, and then feeds the existing packet correspondence endpoint. Thus the Step (x) source equality used by the route can be obtained from realified Frobenioid log-volume compatibility plus the remaining packet source/target calibrations. IUT III, Remark 3.11.4 further says that the root-of-unity and cyclotomic indeterminacies arising from forgetting Frobenius-like labels are invisible for the objects under consideration. The present code formalizes the finite sign-unit part of this assertion: a unit in the subgroup $\{1, -1\} \subset (\mathbb{Z}/\ell\mathbb{Z})^\times$ preserves the zero/nonzero full cusp label, every full-label log-volume expression, and the averaged Gaussian log-volume. This is not yet a construction of the full cyclotomic rigidity isomorphism; it is the finite-label component that the current Stage 1 interface can prove without an additional Frobenioid/cyclotome layer. The same finite-label layer now also proves the corresponding exclusion: an affine coordinate change $j \mapsto t + aj$

preserves the full cusp label if and only if $t = 0$ and a is represented by such a sign-unit invisibility certificate. For nonzero Gaussian environment degree, the same classification is equivalent to point-wise preservation of the Gaussian degree profile. Thus the current code has a finite analogue of the “no further indeterminacies” assertion of Remark 3.11.4(ii). The integrated route now composes the Hodge–Arakelov theta-value sources with this packet-target nonarchimedean source. The resulting theorem derives the same boundary/strict C_Θ dichotomy from source-facing Gaussian evaluation data and source-facing upper-semi data, rather than from two unrelated collections of raw equalities. For the target side, the strongest route now uses a target-only Step (x) alignment

$$\Theta_{\text{avg}} = U_{\text{Ind3}}$$

together with calibration of the local entry target against U_{Ind3} . The direct equality $\Theta_{\text{avg}} = \text{entry}_{\text{target}}$ is then derived. When the theta-source average is identified with the packet-normalized capsule log-volume, this target alignment is itself derived from calibration of U_{Ind3} with the same packet-normalized value. The local entry target may now also be calibrated against this packet-normalized value, so the equality $\text{entry}_{\text{target}} = U_{\text{Ind3}}$ is derived rather than supplied directly in the strongest route.

5.4 Step (xi-f) algebra

The final real algebra is fully formalized. In informal notation, the main theorem is:

Theorem 1 (Formal Step (xi-f)). *Assume $q < 0$, $q \leq \theta$, and $\theta \leq C(-q)$. Then $C \geq -1$. If $q < \theta$, then $C > -1$. If $C = -1$, then $q = \theta$ and $\theta < 0$.*

This theorem is instantiated by the Step (x)-to-Step (xi) route whenever the hull/log-volume upper-ray comparison and the displayed C_Θ -bound are available. It is also instantiated by the current nonarchimedean (Ind3)-plus-factored-SHE route, subject to the same positivity and displayed-bound hypotheses. Thus the remaining gap is not the final ordered-real implication, including its strict and boundary branches, but the construction of the source hypotheses consumed by that implication.

6 Boundary Quotient Experiment

The most developed diagnostic concerns Step (xii), where local Frobenioid p_v^N -ambiguity is compared with the globally calibrated log-volume. The formal model records a local shift

$$L_N = L_0 + N\delta, \quad N \in \mathbb{Z}, \delta \in \mathbb{R},$$

and a global calibration that fixes exponent 0. Lean proves

$$L_N = L_0 \iff N\delta = 0,$$

and, if $\delta \neq 0$,

$$L_N = L_M \iff N = M.$$

At the boundary branch $C_\Theta = -1$, the formal corridor identifies

$$q_{\text{pilot}} = \Theta_{\text{hull}} = L_0.$$

The local shift model is now source-facing: a local Frobenioid log-shell submodule carries a p_v^N -action whose log-volume effect is $L_0 \mapsto L_0 + N\delta$. The local action records whether the p_v^N -endomorphism

is an automorphism, and in the current model this is equivalent to $N = 0$. The global realified Frobenioid calibration package selects exponent 0, hence its calibrated value is L_0 . If the local endomorphism is not an automorphism and $\delta \neq 0$, Lean proves that the local shifted value cannot equal the globally calibrated value. For the actual upper-ray quotient that collapses $\{x \mid x \leq \Theta_{\text{hull}}\}$, Lean proves the exact classification

$$[L_N] = [q_{\text{pilot}}] \iff N\delta \leq 0,$$

and the complementary separation criterion

$$[L_N] \neq [q_{\text{pilot}}] \iff N\delta > 0.$$

If $\delta > 0$, this becomes:

$$[L_N] = [q_{\text{pilot}}] \iff N \leq 0, \quad [L_N] \neq [q_{\text{pilot}}] \iff 0 < N.$$

This is not a proof for either side of the Mochizuki–Scholze–Stix dispute. It is a precise reading of one simplified quotient mechanism inside the current formal corridor. It shows that quotient collapse can identify unequal underlying log-volumes if they lie in the collapsed upper ray, while positive local shifts remain separated. The decisive question is therefore not the ordered-real algebra; it is whether the source IUT constructions actually supply the required calibrated hull/log-volume comparison without making an illicit identification of histories or comparison levels.

7 Current Results

The project currently proves the following source-facing results in Lean:

1. The finite q^{j^2} degree rule, sign quotient, half-range representatives, zero/nonzero split, and average formulas are formalized.
2. The bad-local theta-root quotient- $\mathbb{Z}$ package now feeds the Stage 1 $|F_\ell|$ full-label model through its canonical coordinate; the concrete $Z/\ell Z$ case yields the canonical nonzero full label.
3. A single label-independent scalar cannot absorb all representative j^2 -degrees when $\deg(q) \neq 0$.
4. The raw representative branch does not descend to the sign quotient, while the balanced full-label Gaussian branch does.
5. The Step (x) averaged-log-volume corridor feeds a Step (xi) upper-ray comparison once determinant/hull identifications are supplied.
6. A nonarchimedean (Ind3) entry with factored square/full-label preservation now directly supplies the final weighted-theta route and the signed $C_\Theta \geq -1$ endpoint, once q -positivity and the displayed C_Θ -bound are supplied.
7. The same nonarchimedean route now exposes the strict branch $C_\Theta > -1$ and the boundary consequence $C_\Theta = -1 \Rightarrow q_{\text{signed}} = \Theta_{\text{signed}} < 0$. It also packages these as one dichotomy at the route level.
8. Gaussian degree evaluations with preserved environment degree now feed directly into this nonarchimedean C_Θ dichotomy, constructing the factored square/full-label SHE package internally.

9. In the identity-coordinate Gaussian route, equality at the canonical full label 1 supplies the environment-degree preservation needed by the same nonarchimedean C_Θ dichotomy.
10. The Hodge–Arakelov theta-value evaluation source now constructs the Gaussian degree evaluation from a theta-monoid degree and the bad-local theta-root label source; canonical-label preservation feeds the factored SHE square/full-label obligations.
11. A nonarchimedean upper-semi object equipped with 1-column q-pilot log-Kummer data now supplies the raw nonarchimedean (Ind3) entry alignment used by the canonical Gaussian route.
12. The source-side equality in that compatibility can be derived from product-log-volume preservation along the finite Kummer-plus-forgetting tensor-packet chain.
13. A theta-root attached Kummer/forgetting source package now combines the canonical bad-local $|F_\ell|$ -label source with a source-side packet/product calibration object and that finite chain, deriving the nonarchimedean source equality.
14. Source-side and target-side packet calibrations are now joined by one nonarchimedean log-Kummer packet-correspondence source carrying the q-pilot non-interference guard. Its projections derive the theta-to-packet, theta-to-entry, entry-to-(Ind3)-target, and packet-normalized target equalities used by the route; the target calibration now refers directly to the packet-normalized capsule value rather than to a separate real representative.
15. The strongest Gaussian first-pass route can now consume this single packet-correspondence source and project both its source and target faces, instead of taking source packet equalities and target calibration as separate route inputs.
16. A realified-Frobenioid log-Kummer entry source now bundles the selected upper-semi inclusion with the realified packet source; the strongest Step (x) route projects Kummer transfer, mono-analytic forgetting, source/target packet calibration, and q-pilot non-interference from this single source.
17. A theta-root attached realified entry source now combines the canonical bad-local theta-root label package with that realified Step (x) source, and the Hodge–Arakelov route can consume this attached source directly.
18. The Hodge/SHE/IPL/hull route now also has a realified endpoint: route alignment supplies the Hodge–Arakelov and q-pilot data, and Step (x) enters through the theta-root attached realified source.
19. Remark 3.9.5(iii)’s hull-approximant family is now represented directly: an approximant lies inside the canonical hull $\phi(P)$, has log-volume between P and $\phi(P)$, and is fixed by hull formation.
20. The bounded-family upper-semi quotient of Remark 3.9.5(v)–(vi) is represented: each possible region lies in the hull of the family union, and any two nonempty possible regions have identical images in the quotient that collapses this family hull.
21. The Corollary 3.12 possible-image approximant source now projects to this bounded-family quotient source, yielding the same quotient-image equality for nonempty possible theta-image choices.

22. In that same source, q -pilot containment in the chosen approximant implies q -pilot containment in the family hull; if the q -pilot region is nonempty, its quotient image agrees with the quotient image of any nonempty possible theta-image choice.
23. Such a certified hull-approximant now also constructs the Step (xi) upper-ray datum: if the q -pilot region is contained in the approximant and the approximant log-volume is determinant-normalized, then Lean derives the q -pilot bound by the approximant, by the canonical hull, and by the determinant log-volume.
24. The approximant route is now specialized to the Corollary 3.12 possible-image union. A single source object binds the possible theta-image family, its union, the chosen hull-approximant, q -pilot containment in that approximant, and determinant normalization; the experiment endpoint returns the union/approximant/hull and q -pilot/determinant inequalities.
25. The determinant input can now be produced from a finite weighted localization source: adjusted vector-bundle summand log-volumes are summed, the positive tensor-power value is derived, and normalization is proved to return the determinant log-volume before projecting to the older determinant interface.
26. Each localization summand can now itself be produced from a finite direct-summand source: the bundle log-volume is the sum of direct-summand log-volumes, then the structure-sheaf adjustment and positive weight produce the adjusted localization contribution.
27. The Ob3-1-2 structure-sheaf adjusted source makes the $\mathcal{O}(-)$ subtraction explicit before projection to the localization summand. A corresponding determinant constructor proves that summing these weighted adjusted sources gives the determinant log-volume used by the route.
28. The possible-image hull-approximant endpoint now has a weighted determinant-source constructor. The route consumes the weighted summand source and projects to the common upper-ray interface only after a named Ob3-3 compatibility certificate identifies the approximant log-volume with the weighted normalized determinant value.
29. The Ob4 tensor-power passage is now represented at log-volume level: M -th tensor-power formation multiplies by $M > 0$, and normalization recovers the base log-volume. The weighted determinant source projects to this tensor-power object, and two tensor-power presentations with the same base log-volume have equal normalized log-volume.
30. A first realified Frobenioid degree layer is present: equality of realified Frobenioid log-volumes now derives the product-log-volume equality used by the tensor-packet coric transfer, with theater histories still kept distinct. The degree object also has a finite tensor-product operation: divisor degree, unit log-volume, and realified log-volume are proved to add. A tensor-packet source certified by this Frobenioid product inherits product-log-volume additivity.
31. The global-to-local realified Frobenioid restriction factor of IUT I, Example 3.5 is formalized as inverse-extension-degree scaling of the local prime log-volume, with the degree-normalization identity proved in Lean. This restricted value is connected to the local p_v^N -shift source used in the Step (xii) diagnostic.
32. The diagonal/arbitrary distribution distinction of IUT I, Example 5.1 and Remark 5.1.1 is formalized at finite realified log-volume level: diagonal distributions are constant in j , their total is $|J|$ times the common restricted value, and a two-coordinate mismatch cannot be represented as diagonal.

33. The pivotal distribution of IUT I, Example 5.4(vii), is modeled by a chosen index in a diagonal distribution; Lean proves the summed log-volume formula and the normalized averaged formula of Remark 5.4.2.
34. The finite NF-bridge log-volume model of IUT I, Definition 5.5, separates summed and averaged bridge targets and proves their $|J|$ scaling relation for a common capsule distribution.
35. The finite Θ NF-bridge compatibility model records the same-index-bijection and same-capsule requirements of Definition 5.5(iii), and proves the corresponding Θ -target equalities for summed and averaged NF modes.
36. The finite gluing set of Corollary 5.6(iii) is modeled as the additive F_ℓ -torsor on $\mathbb{Z}/\ell\mathbb{Z}$, with compatibility data preserved under gluing translation.
37. The realified Frobenioid layer now feeds the nonarchimedean log-Kummer packet source: Kummer and mono-analytic forgetting transfers are projected from Frobenioid compatibility before the packet correspondence endpoint is applied.
38. A packet-normalized Step (x) entry-target source now constructs the target side of a nonarchimedean upper-semi entry from the capsule-family normalized log-volume. The target-calibration constructor derived from this source removes the separate entry-target-to-packet-normalized equality.
39. The same target-calibration layer can now derive the (Ind3)-target-to-packet-normalized equality from theta-target alignment and theta-average packet normalization, instead of accepting that equality directly.
40. The log-Kummer packet-correspondence source can now be constructed from source packet/product calibration plus the packet-normalized entry target and theta-target alignment. The q-pilot non-interference, source equality, and target equality endpoint is then recovered from this constructed correspondence.
41. The source packet/product calibration itself now has a constructor from packet-local/source alignment, packet-local (Ind3)-source alignment, and Kummer/forgetting product-log-volume preservation. The source holomorphic-product equality is derived by this constructor.
42. The packet-correspondence constructor has a stronger alignment form: packet-local source alignment, packet-local (Ind3)-source alignment, packet-normalized entry target, and theta-target alignment generate both packet calibrations before projecting the q-pilot endpoint.
43. The realified-Frobenioid log-Kummer packet source has the same reduced target-data constructor: Frobenioid compatibility yields the Kummer and forgetting transfers, while packet-normalized target data is generated from the entry-target source and theta-target alignment.
44. The realified packet source now also has a source-reduced constructor: packet-local/source alignment and packet-local (Ind3)-source alignment generate the source packet calibration using the Frobenioid-derived Kummer and forgetting transfers.
45. The realified Step (x) entry source has the same reduced boundary: entry membership plus source/target alignment data constructs the packet source internally and projects q-pilot non-interference, source equality, target equality, and the upper-semi inequality.

46. The same realified packet source now exposes a combined calibration endpoint: it proves Kummer and forgetting product-log-volume preservation, packet-local equality with the entry source, entry-source equality with the (Ind3) source, packet-normalized equality for the theta average, (Ind3) target, and entry target, and q-pilot non-interference.
47. The theta-root attached realified entry source now has the same reduced constructor and projects both the canonical nonzero theta-root label facts and the realified Step (x) upper-semi endpoint.
48. The Hodge/SHE/IPL/hull route now consumes this reduced theta-root Step (x) data directly: the theorem constructs the realified source internally and concludes the current C_Θ dichotomy.
49. The finite sign-unit/root-of-unity part of Remark 3.11.4 is formalized: sign-unit action preserves full cusp labels, full-label log-volume readings, Gaussian degrees, and averaged Gaussian log-volume.
50. The same finite layer classifies the allowed unit-affine label changes: full-label preservation is equivalent to zero translation plus a sign-unit invisibility certificate; under nonzero Gaussian environment degree this is also equivalent to pointwise Gaussian-degree preservation.
51. The current integrated route composes Hodge–Arakelov theta-value sources with the packet-target nonarchimedean source to obtain the same $C_\Theta = -1$ boundary alternative or strict $C_\Theta > -1$ conclusion.
52. A Theorem 3.11 multiradial source record now packages structured IPL/SHE/APT, coric output, possible theta images, pilot-column asymmetry, distinct log-shell treatments, and the no-history-identification guard.
53. The IPL source layer now has an explicit log-volume transport record with source and target Hodge-theater sides, prime-strip endpoint compatibility, log-volume preservation, and no-history-identification.
54. The SHE source layer now derives the current square/full-label preservation package from Hodge–Arakelov canonical-one synchronization, exposing pointwise representative comparison, transported-average equality, and the retained Hodge-history separation.
55. The Step (xi) hull/determinant layer now has a Theorem 3.11 source constructor: source obligations are derived from the source record where possible, the record’s theta possible-image union is tied to the hull endpoint, and the determinant bound yields $q_{\text{signed}} \leq \theta_{\text{signed}}$.
56. The package-level hull/determinant source bridge can now be constructed from a concrete common hull containing the target union and a determinant/log-volume bound for that hull. The experiment endpoint returns target-union containment, the determinant bound, all-target bounds, and equality with the package’s named hull/determinant bridge.
57. The Step (x)-to-Step (xi) upper-ray bridge can now be built directly from a certified possible-image hull approximant. The determinant normalization is read from the approximant/determinant compatibility, rather than supplied as a separate equality at the bridge boundary.
58. A Hodge/SHE/IPL/hull packet-target route now derives the target-side log-volume equality from SHE synchronization and reads q -positivity from the hull constructor; the source-side

theta/Hodge equality is derived from a full-label calibration certificate in the route-alignment source record.

59. The target-side theta-to-entry equality is no longer a direct input in the strongest route. It is derived from a Step (x) theta-target alignment and calibration of the entry target with the upper-semi target value.
60. The Step (x) theta-target alignment can itself be derived from packet-normalized theta-source equality and calibration of the upper-semi target with the packet-normalized value.
61. The entry-target calibration can likewise be expressed by equality to the packet-normalized value, deriving equality with the upper-semi target.
62. The Step (xi-f) real algebra proves $C_\Theta \geq -1$, strictness under strict upper-ray membership, and equality constraints in the boundary case.
63. The final route rejects the balanced sign-compatible comparison level as sufficient; it requires hull/log-volume comparison data.
64. The local Frobenioid shift model is now sourced by a p_v^N log-volume action package and a global realified calibration package; a non-automorphic local shift with nonzero step is separated from the calibrated value, and the upper-ray quotient experiment gives the exact threshold for collapse versus separation in the boundary upper ray.
65. Several downstream IUT IV numerical estimates are formalized as ordered-real algebra conditional on the Corollary 3.12 input.

The first-pass experiment report evaluates the route-theorem, finite Gaussian, collapse-obstruction, upper-ray, and IUT IV algebra flags to **true**. The flag **disputeSettledByCurrentStage** remains **false**.

The current code base is large enough to distinguish three meanings of “available”. First, pure ordered-real consequences are proved outright. Second, route theorems prove that certain named source-facing records imply the desired endpoint. Third, the records themselves remain obligations when they stand for Hodge-theater, Frobenioid, log-Kummer, hull, determinant, or multiradial constructions not yet formalized from IUT I–III. Only the first two meanings are currently achieved.

8 What Is Not Yet Proved

No formal statement currently proves IUT III, Corollary 3.12 from Mochizuki’s published hypotheses. The following mathematics is still external to the code:

1. initial Θ -data and its arithmetic hypotheses from IUT I;
2. Hodge theaters, Frobenioids, log-shells, prime strips, and theta links;
3. the Hodge–Arakelov theta evaluation and Gaussian monoids of IUT II;
4. the construction of the actual (Ind1), (Ind2), (Ind3) data in Theorem 3.11;
5. the genuine holomorphic hull and determinant constructions used in Step (xi), beyond the present source-data constructor and ordered-real upper-ray model;

6. the full proof that the typed SHE and Hodge-descent transport obligations follow from Mochizuki’s source constructions, beyond the current Hodge–Arakelov canonical-one synchronization shadow;
7. the construction of the input prime-strip link (IPL), simultaneous holomorphic expressibility (SHE), and algorithmic parallel transport (APT) properties as actual consequences of Theorem 3.11 rather than records;
8. the arithmetic-divisor, height, log-different, and log-conductor constructions behind the IUT IV Diophantine applications.

9 Implementation Audit

The code contains no `sorry`, `admit`, or Lean axioms in the Stage 1 corridor. The current risk is not hidden unsoundness of this form, but mathematical incompleteness: several records are deliberately named obligation interfaces. These interfaces are acceptable only as temporary boundaries. The next phase must replace them by constructions from the source papers or mark them as unresolved.

The endpoint type `Stage1Comparison` is not a source proof: its comparison field is the Corollary-3.12-shaped conclusion, and the corresponding theorem is only a projection. The source-facing route must construct this endpoint from the ledger. The ledger now records the opaque output certification and ties the Θ -common bound used in the final inequality to `chartedContainer.apply_certificate`, so the displayed charted container cannot be paired with an unrelated convenient bound. Throughout the current corridor, “upper ray” means the signed-log region $\{x \mid x \leq b\}$, not the ordinary order-theoretic ray $\{x \mid b \leq x\}$.


The largest engineering issue is `Iut/Stage1/IUTStage1Source.lean`, which currently contains most of the finite-label, Gaussian, Hodge/SHE, hull, Step (x), Step (xi), Step (xii), and IUT IV algebra. This is now too large for comfortable review. Low-risk future refactors are:

finite labels and q^{j^2} averages	separate module,
Step (xi) hull/determinant upper rays	separate module,
Step (xii) local Frobenioid shifts	separate module,
IUT IV real algebra	separate module.

These refactors should be mechanical only: move declarations, preserve names where possible, and rebuild after each move.

10 Next Work

The next mathematical milestone is to construct, rather than assume, more of the source-facing SHE and Hodge-descent preservation data. The immediate route is:



The project should continue by replacing the present record fields for factored SHE preservation and Hodge-descent source-target alignment with the corresponding constructions from IUT I–III.

The practical priority order is:

1. factor the large source module into reviewable finite-label, hull, Frobenioid-shift, and IUT IV algebra modules without changing theorem statements;
2. extend the Gaussian-derived factored SHE construction beyond the finite degree-evaluation and canonical-label shadows toward actual Gaussian/Frobenioid material from IUT II;
3. construct the packet-correspondence source from the actual log-Kummer/Frobenioid machinery of IUT III, Step (x), rather than treating this finite correspondence object itself as source-facing input;
4. replace the hull/determinant endpoint records by the actual holomorphic-hull and determinant operations from Remark 3.9.5 and Step (xi);
5. only then evaluate whether the Stage 1 interface has enough source content to bear on the Mochizuki–Scholze–Stix disagreement.

References

- [1] S. Mochizuki, *Inter-universal Teichmüller Theory I: Construction of Hodge Theaters*.
- [2] S. Mochizuki, *Inter-universal Teichmüller Theory II: Hodge-Arakelov-theoretic Evaluation*.
- [3] S. Mochizuki, *Inter-universal Teichmüller Theory III: Canonical Splittings of the Log-theta-lattice*.
- [4] S. Mochizuki, *Inter-universal Teichmüller Theory IV: Log-volume Computations and Set-theoretic Foundations*.
- [5] S. Mochizuki, *On the Formalization of IUT: Preliminary Progress Report*, 2026.
- [6] P. Scholze and J. Stix, *Why abc is still a conjecture*.